



MarioGPT: Open-Ended Text2Level Generation through Large Language Models

Viet Nhat Nguyen (21520378)

Thi Lien Le (21522282)

July 6, 2023



UNIVERSITY OF
INFORMATION
TECHNOLOGY



Table of Contents

1 Introduction

- ▶ Introduction
- ▶ Background and Related Work
- ▶ Open-Ended Level Generation through LLMs
- ▶ Experiments and Results
- ▶ Conclusion

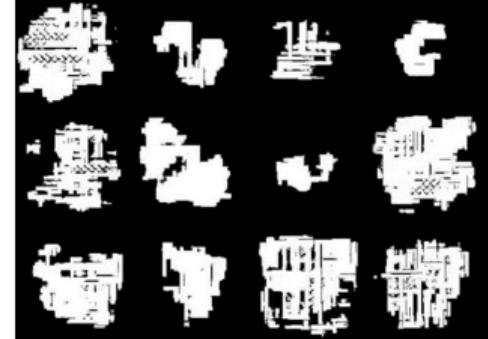


PCG and Machine Learning

1 Introduction



Doom maps generated by
GANs model





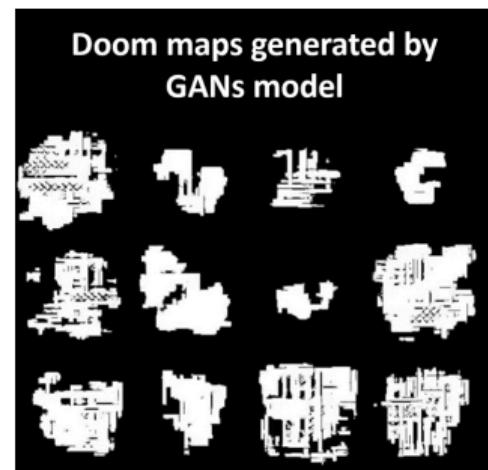
PCG and Machine Learning

1 Introduction

- PCG for evaluating generalization capabilities of trained agents.
 - Offer more diverse and challenging scenarios.
 - A better way to test the adaption and generalization of trained agents.



Arcade Learning Environment



Doom maps generated by
GANs model



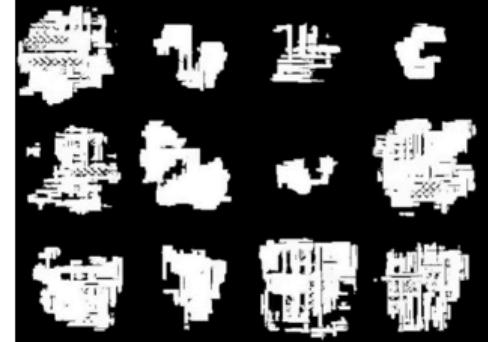
PCG and Machine Learning

1 Introduction

- PCG for evaluating generalization capabilities of trained agents.
 - Offer more diverse and challenging scenarios.
 - A better way to test the adaption and generalization of trained agents.
- Incorporation of machine learning-based approaches in PCG systems.
 - GANs for generating levels (Doom, Super Mario Bros)



Doom maps generated by
GANs model





Challenges

1 Introduction

- **Costly** searching in the latent space of neural networks.
- Desire to **directly condition** a generator for creating levels with specific properties, ideally in natural language. (**prompt**)



Contributions

1 Introduction

- Introduction of MarioGPT, a text-to-level model generating Mario levels based on natural language prompts.



Contributions

1 Introduction

- Introduction of MarioGPT, a text-to-level model generating Mario levels based on natural language prompts.
- Combining MarioGPT with novelty search for producing diverse levels.



Table of Contents

2 Background and Related Work

- ▶ Introduction
- ▶ Background and Related Work
- ▶ Open-Ended Level Generation through LLMs
- ▶ Experiments and Results
- ▶ Conclusion



Procedural Content Generation (PCG)

2 Background and Related Work

- Procedural Content Generation (PCG) refers to techniques that can automatically create game content (e.g. levels, maps, or characters).
 - Increasing the replayability
 - Reducing production costs



Procedural Content Generation (PCG)

2 Background and Related Work

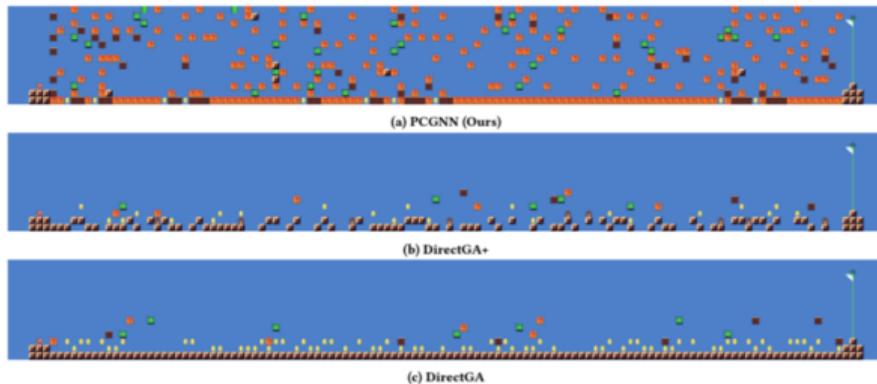


a) Rouge b) Elite c) Diablo III d) Minecraft e) No Man's Sky f) Civilization VI

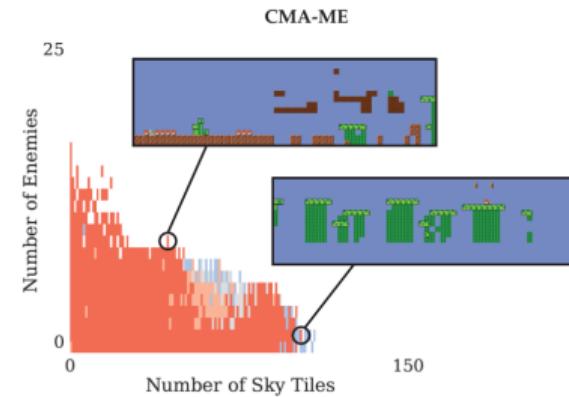


Neural Network-based Level Generation

2 Background and Related Work



Using neural networks ([Beukman et al.\[1\]](#))



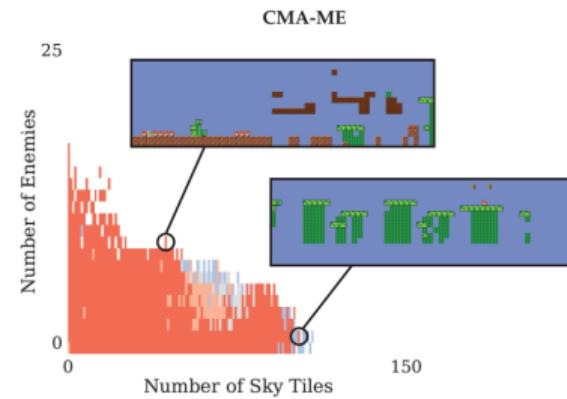
Searching in the latent space of GANs
([Matthew et al.\[2\]](#))



Neural Network-based Level Generation

2 Background and Related Work

- Guided sampling of the latent space could result in a diverse set of levels.
- The abilities to control characteristics in generated levels.



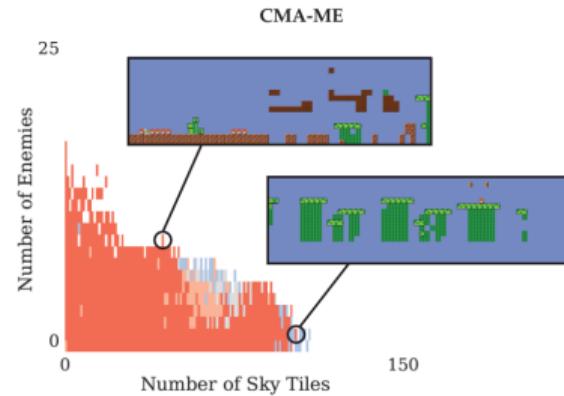
Searching in the latent space of GANs
([Matthew et al.\[2\]](#))



Neural Network-based Level Generation

2 Background and Related Work

- Guided sampling of the latent space could result in a diverse set of levels.
- The abilities to control characteristics in generated levels.
- *Using A* agent to measure whether the generated levels were playable.*



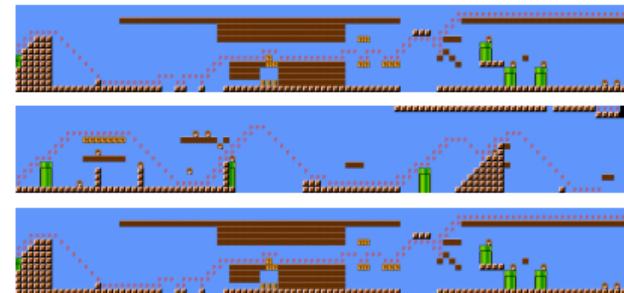
Searching in the latent space of GANs
([Matthew et al.\[2\]](#))



Open-Endedness Paradigm

2 Background and Related Work

- The open-endedness paradigm focuses on algorithms that can produce infinite innovation.



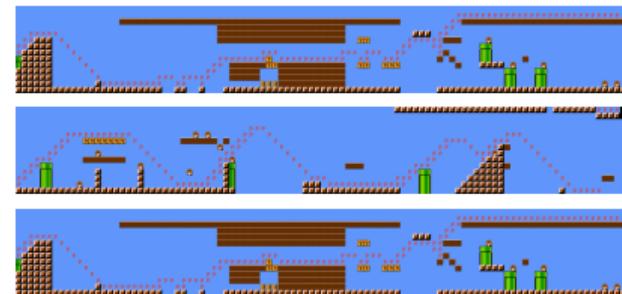
Diversity but not-playable cases



Open-Endedness Paradigm

2 Background and Related Work

- The open-endedness paradigm focuses on algorithms that can produce infinite innovation.
- Must balance the hard task of generating content with **diversity** as well as **playability**



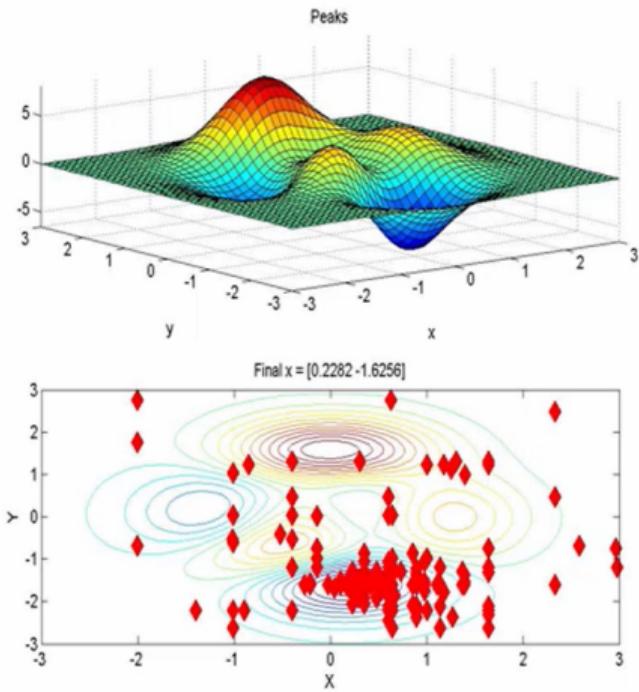
Diversity but not-playable cases



Genetic Algorithms

2 Background and Related Work

- Uses concepts from *evolutionary biology*

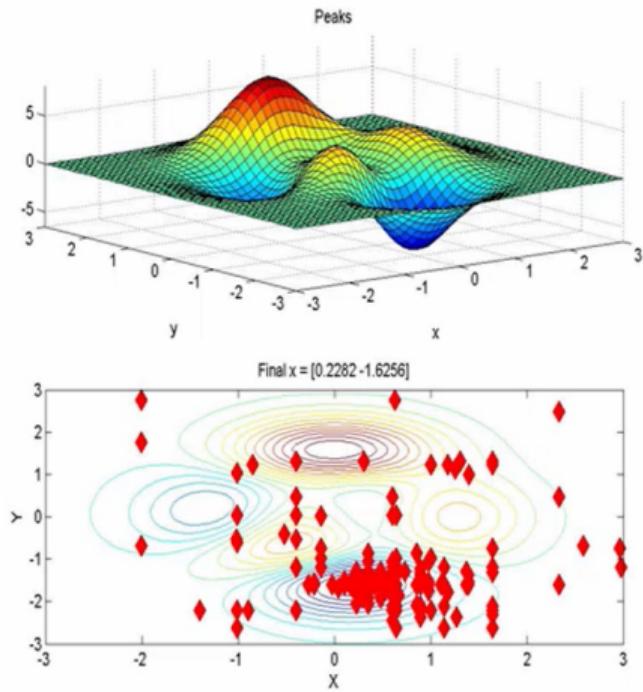




Genetic Algorithms

2 Background and Related Work

- Uses concepts from *evolutionary biology*
- Start with an initial generation of candidate solutions that are tested against the objective function

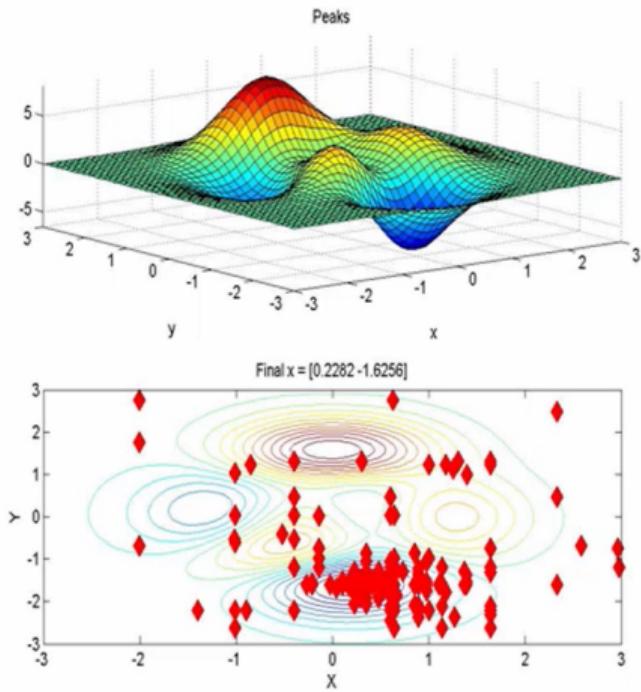




Genetic Algorithms

2 Background and Related Work

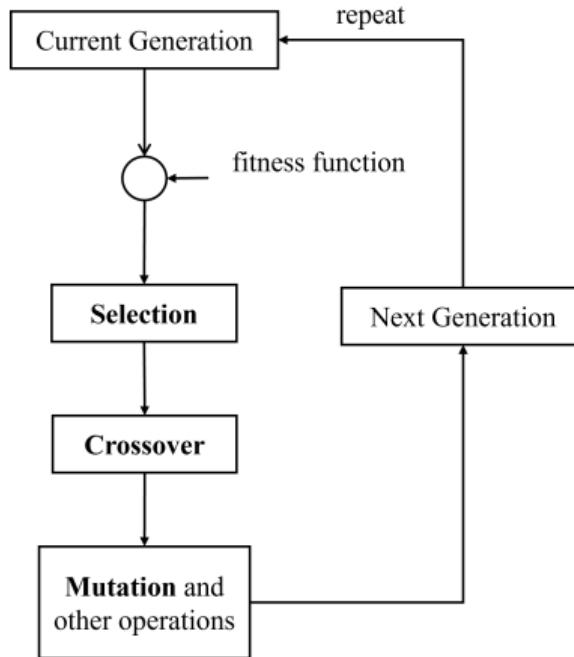
- Uses concepts from *evolutionary biology*
- Start with an initial generation of candidate solutions that are tested against the objective function
- Subsequent generations evolve from the 1st through *selection, crossover* and *mutation*





Genetic Algorithms

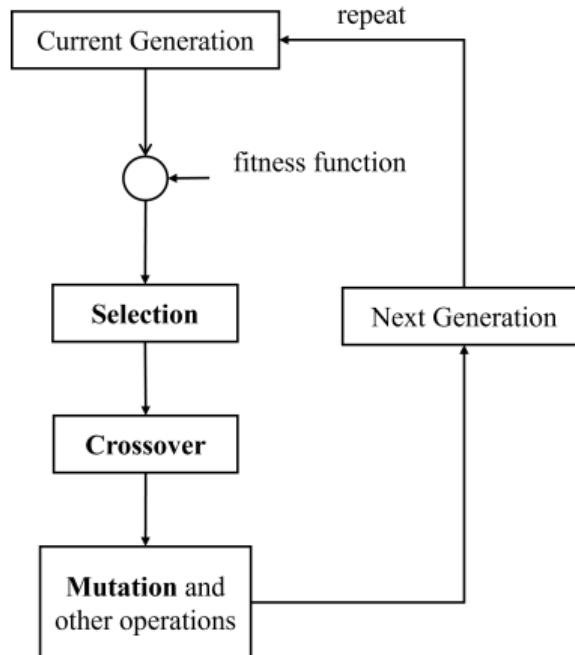
2 Background and Related Work





Genetic Algorithms

2 Background and Related Work

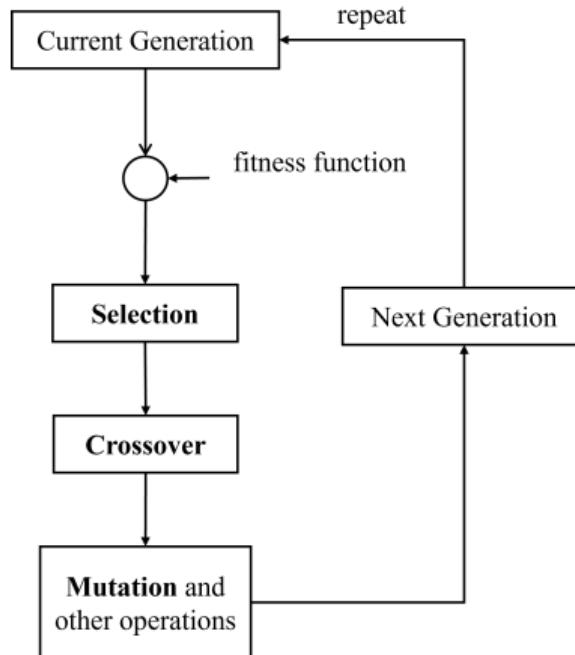


- First generation is arbitrarily, or $x = \text{random}()$



Genetic Algorithms

2 Background and Related Work

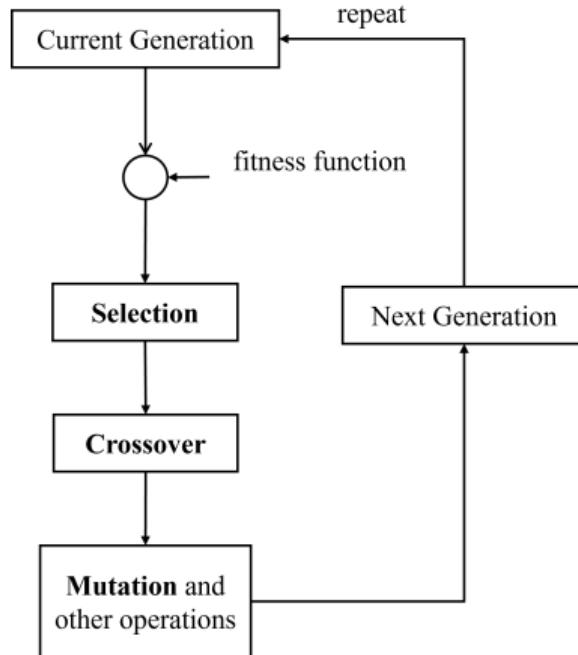


- First generation is arbitrarily, or $x = \text{random}()$
- Evaluate their *fitness*



Genetic Algorithms

2 Background and Related Work



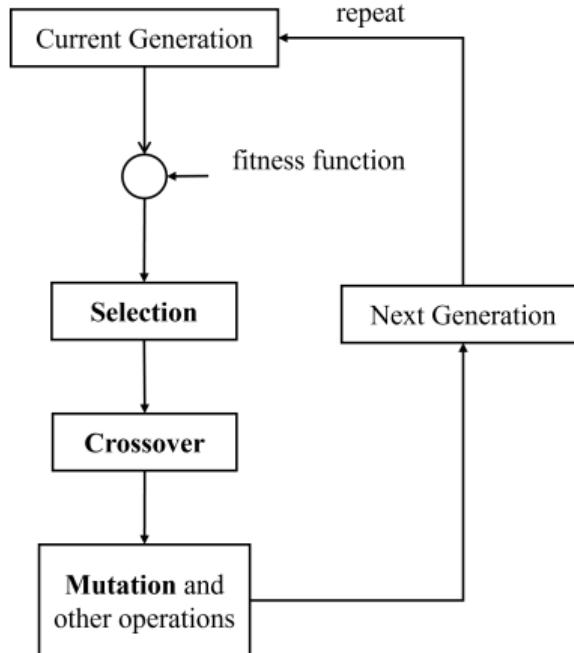
- First generation is arbitrarily, or $x = \text{random}()$
- Evaluate their *fitness*

$$\text{fitness}(x) = f(x)$$



Genetic Algorithms

2 Background and Related Work

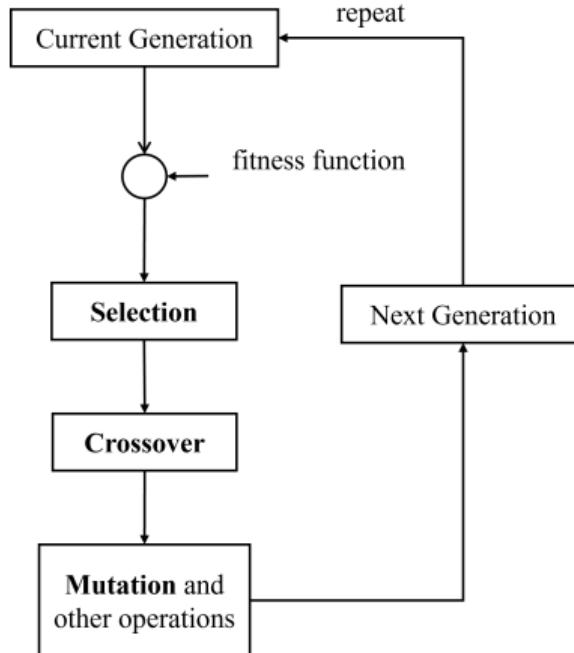


- First generation is arbitrarily, or $x = \text{random}()$
- Evaluate their *fitness*
- *Select* members of the population:



Genetic Algorithms

2 Background and Related Work

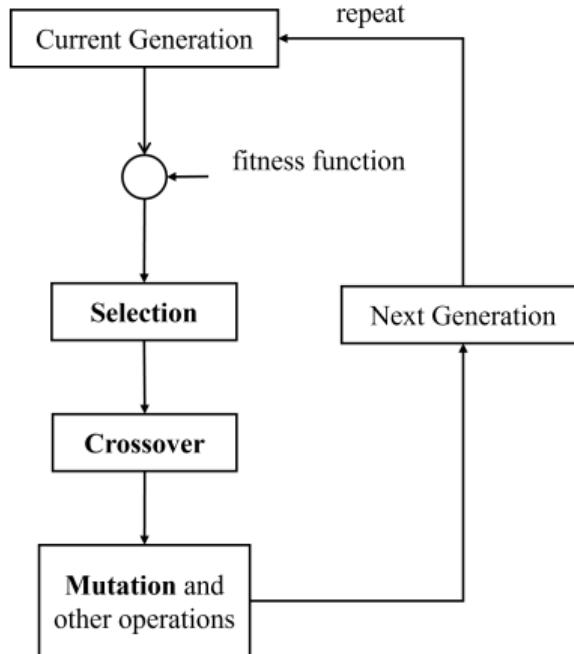


- First generation is arbitrarily, or $x = \text{random}()$
- Evaluate their *fitness*
- Select members of the population:
 - Deterministic



Genetic Algorithms

2 Background and Related Work

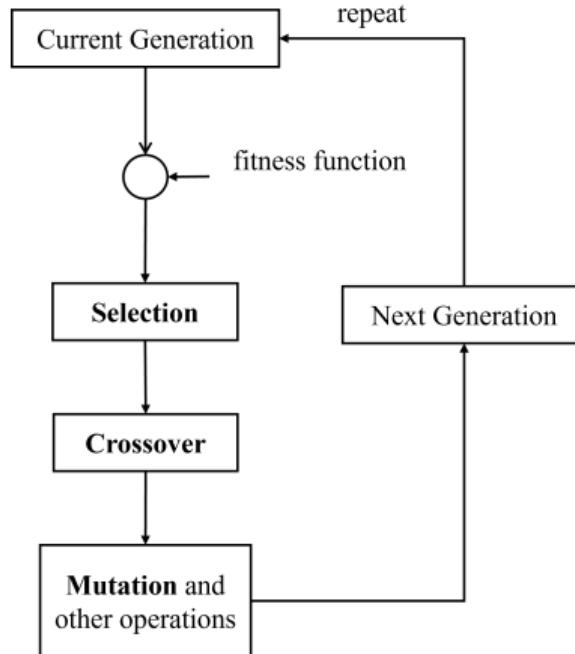


- First generation is arbitrarily, or $x = \text{random}()$
- Evaluate their *fitness*
- Select members of the population:
 - Deterministic: Elitism
 $\text{argmax}_x \text{fitness}(x)$



Genetic Algorithms

2 Background and Related Work

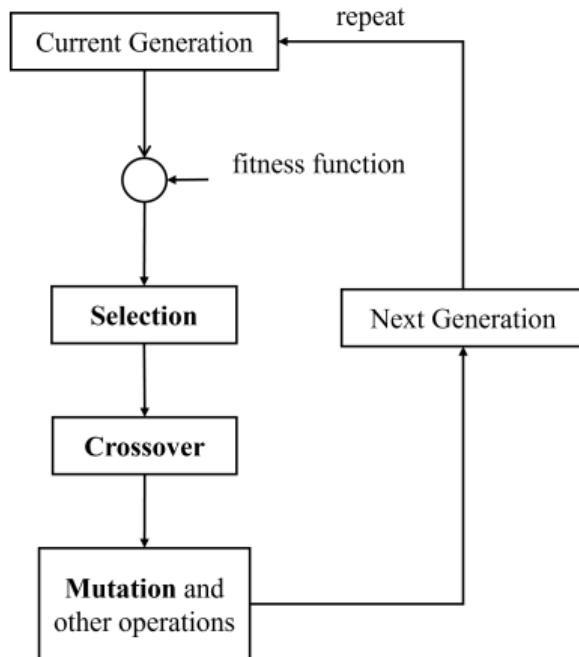


- First generation is arbitrarily, or $x = \text{random}()$
- Evaluate their *fitness*
- Select members of the population:
 - Deterministic
 - Stochastic



Genetic Algorithms

2 Background and Related Work



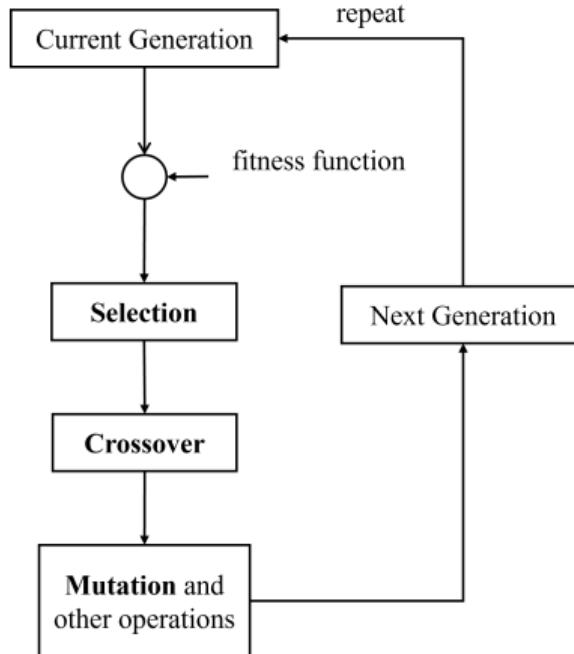
- First generation is arbitrarily, or $x = \text{random}()$
- Evaluate their *fitness*
- Select members of the population:
 - Deterministic
 - Stochastic: Roulette wheel selection

$$P(x) = \frac{\text{fitness}(x)}{\sum \text{fitness}(x)}$$



Genetic Algorithms

2 Background and Related Work

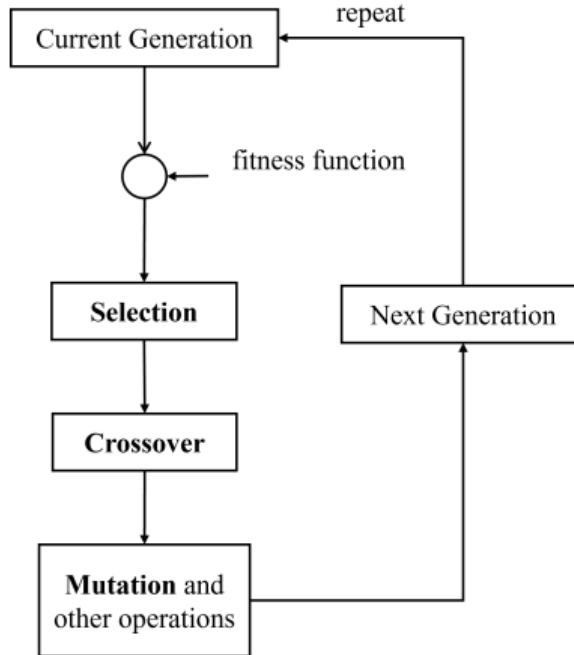


- First generation is arbitrarily, or $x = \text{random}()$
- Evaluate their *fitness*
- Select members of the population:
 - Deterministic
 - Stochastic



Genetic Algorithms

2 Background and Related Work

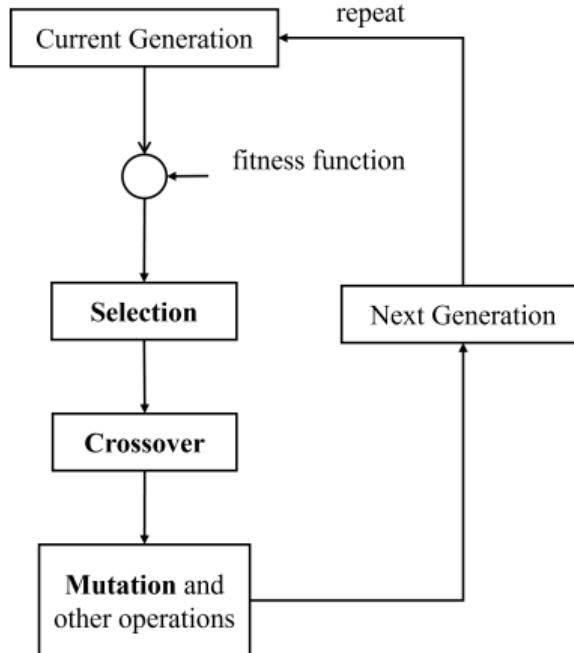


- First generation is arbitrarily, or $x = \text{random}()$
- Evaluate their *fitness*
- Select members of the population:
 - Novelty Selection



Genetic Algorithms

2 Background and Related Work

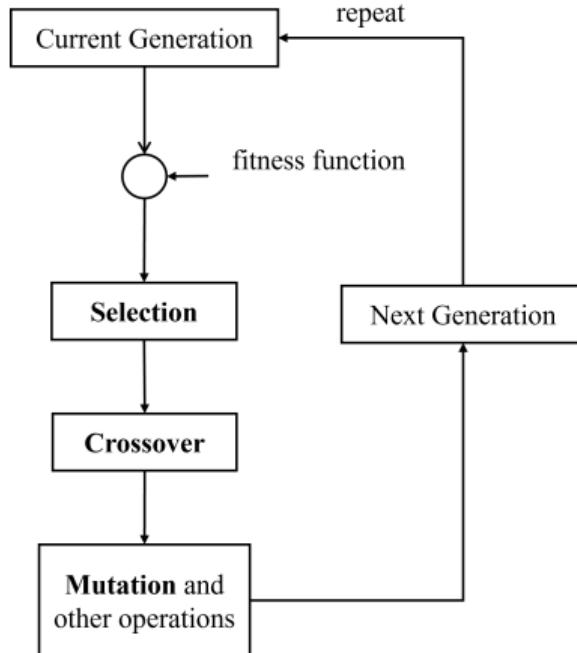


- First generation is arbitrarily, or $x = \text{random}()$
- Evaluate their *fitness*
- Select members of the population:
 - Novelty Selection



Genetic Algorithms

2 Background and Related Work

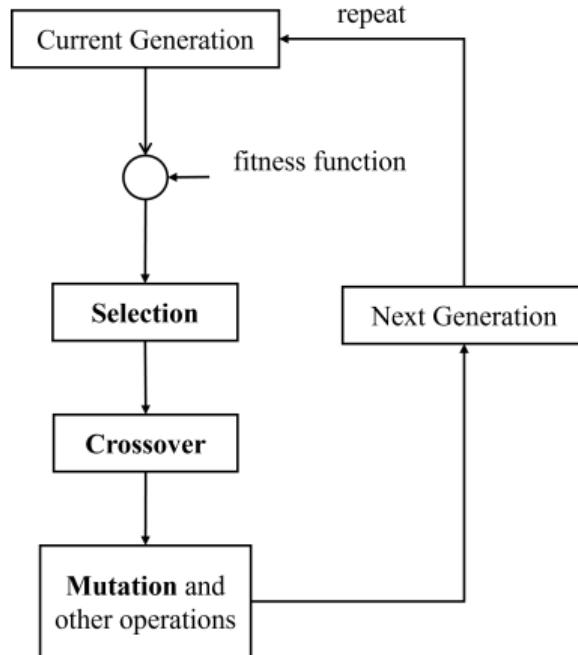


- First generation is arbitrarily, or $x = \text{random}()$
- Evaluate their *fitness*
- Select members of the population:
 - Novelty Selection
- Implement crossover operation on the reproduced chromosomes



Genetic Algorithms

2 Background and Related Work



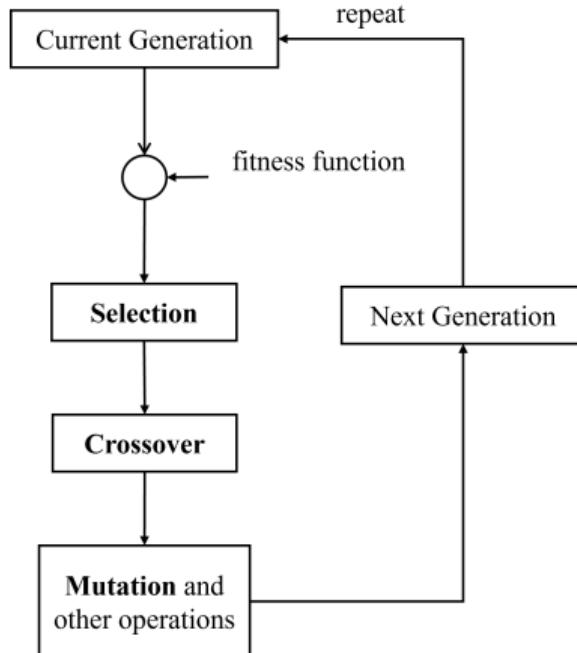
- First generation is arbitrarily, or $x = \text{random}()$
- Evaluate their *fitness*
- Select members of the population:
 - Novelty Selection
- Implement crossover operation on the reproduced chromosomes

$$\text{Crossover}(x_1, x_2) = y$$



Genetic Algorithms

2 Background and Related Work

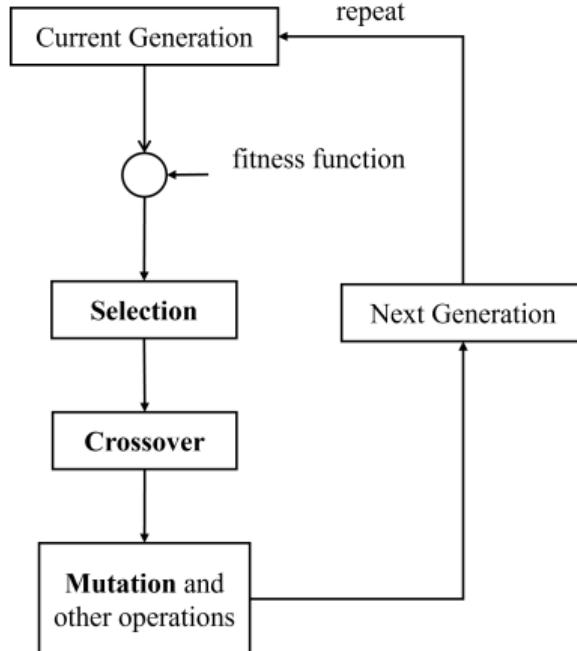


- First generation is arbitrarily, or $x = \text{random}()$
- Evaluate their *fitness*
- Select members of the population:
 - Novelty Selection
- Implement crossover operation on the reproduced chromosomes
$$\text{Crossover}(x_1, x_2) = y$$
- Execute *mutation* operation with low probability



Genetic Algorithms

2 Background and Related Work

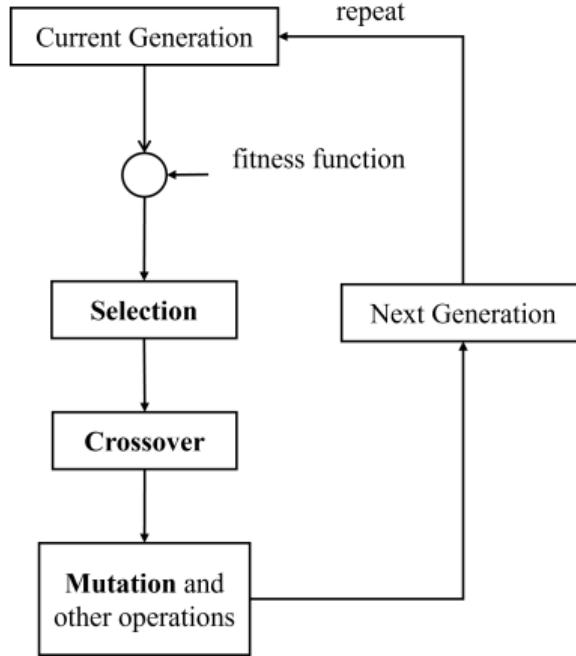


- First generation is arbitrarily, or $x = \text{random}()$
- Evaluate their *fitness*
- Select members of the population:
 - Novelty Selection
- Implement *crossover* operation on the reproduced chromosomes
$$\text{Crossover}(x_1, x_2) = y$$
- Execute *mutation* operation with low probability
$$\text{Mutation}(y) = y', \text{ with } p_m$$



Genetic Algorithms

2 Background and Related Work



- First generation is arbitrarily, or $x = \text{random}()$
- Evaluate their *fitness*
- Select members of the population:
 - Novelty Selection
- Implement *crossover* operation on the reproduced chromosomes
$$\text{Crossover}(x_1, x_2) = y$$
- Execute *mutation* operation with low probability
$$\text{Mutation}(y) = y', \text{ with } p_m$$
- Allow the integration of multiple objectives



Sequence Modelling and Transformer

2 Background and Related Work

- Classic approaches to sequence modelling are RNNs ([Rumelhart et al.](#)[4]) and LSTM networks ([Hochreiter et al.](#) [3])
 - fading memory
 - limited scalability
- Transformers ([Vaswani et al.](#)[5]) can address both challenges, can be parallelized and achieve good scalability!



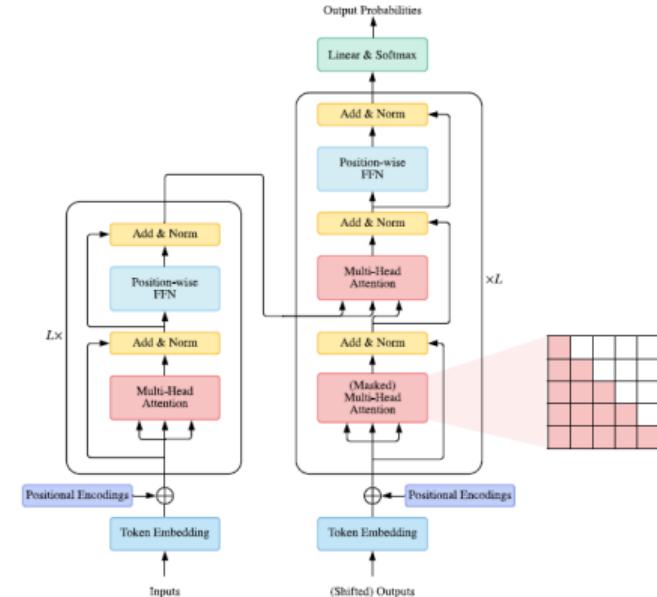
Sequence Modelling and Transformers

2 Background and Related Work

- "Attention is all you need" ([Vaswani et al., 2017](#))

- Encoder:

- Decoder:

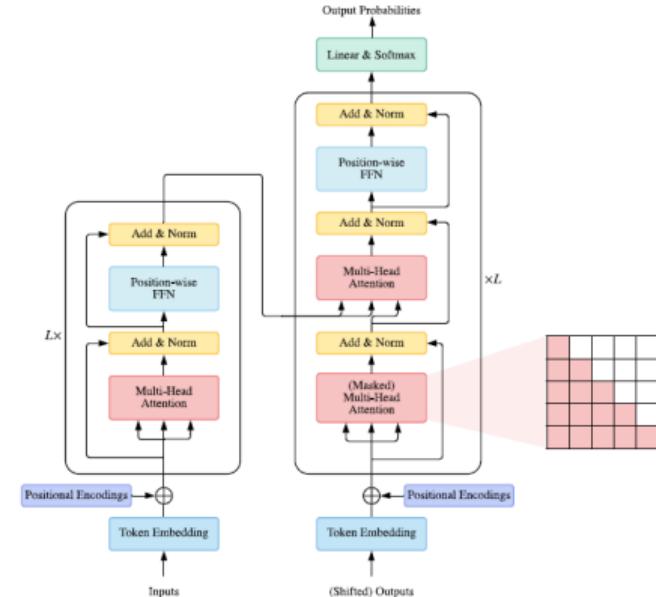




Sequence Modelling and Transformers

2 Background and Related Work

- "Attention is all you need" ([Vaswani et al., 2017](#))
- **Encoder:**
 - N identical layers (Multi-Head Attention and FFN)
 - Residual connections
- **Decoder:**

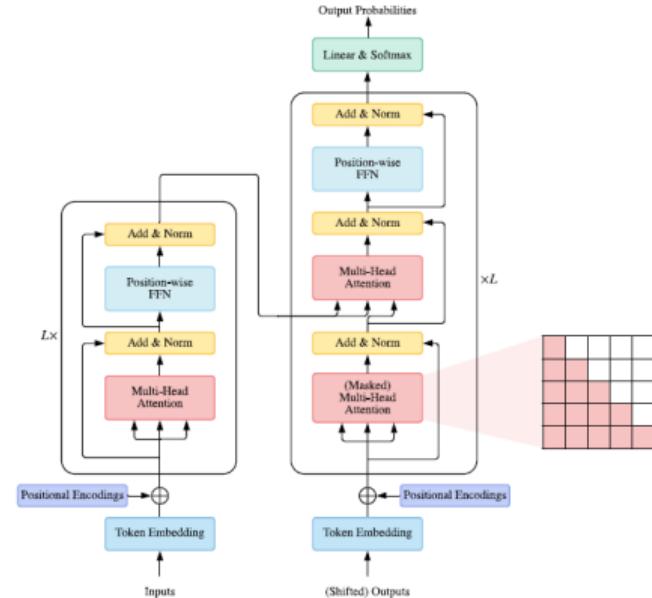




Sequence Modelling and Transformers

2 Background and Related Work

- "Attention is all you need" ([Vaswani et al., 2017](#))
- **Encoder:**
 - N identical layers (Multi-Head Attention and FFN)
 - Residual connections
- **Decoder:**
 - N identical layers (Multi-Head Attention, FFN and Masked Multi-Head Attention)
 - Residual connections

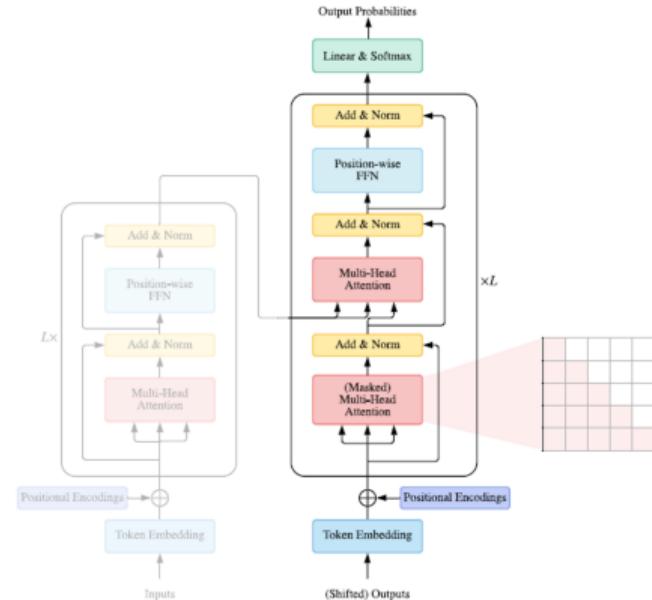




Sequence Modelling and Transformers

2 Background and Related Work

- "Attention is all you need" ([Vaswani et al., 2017](#))
- **Encoder:**
 - N identical layers (Multi-Head Attention and FFN)
 - Residual connections
- **Decoder:**
 - N identical layers (Multi-Head Attention, FFN and Masked Multi-Head Attention)
 - Residual connections
- *GPT consists of a stack of decoder layers.*

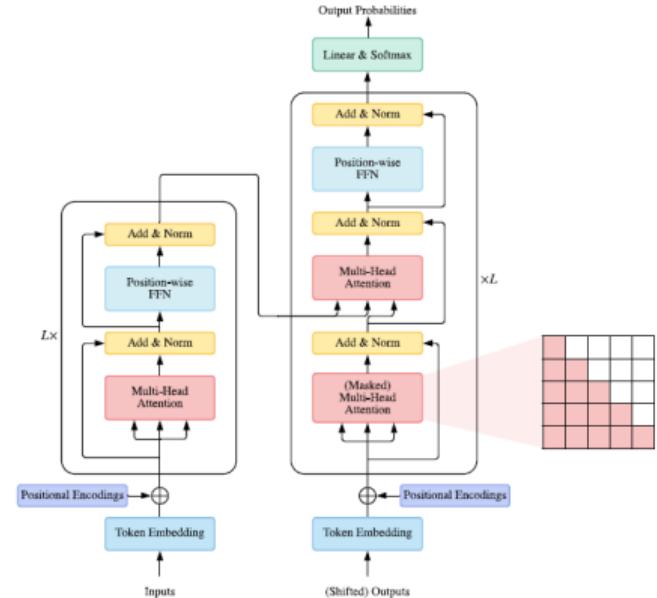




Sequence Modelling and Transformers

2 Background and Related Work

- Enable LLMs to learn from massive datasets

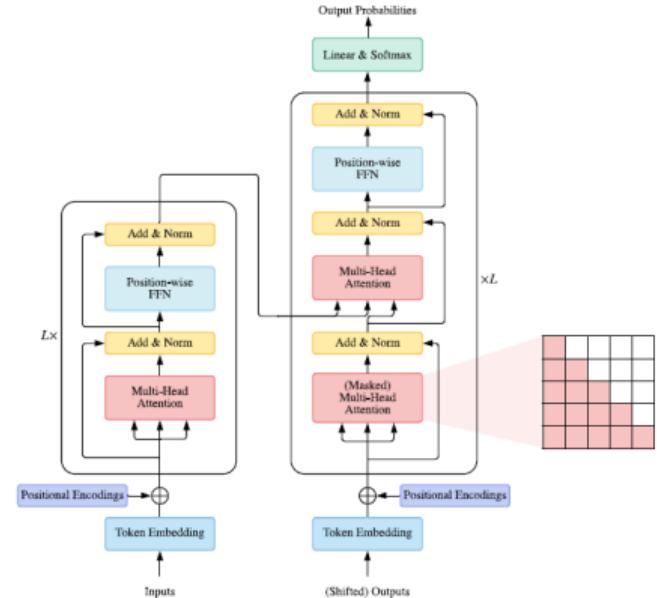




Sequence Modelling and Transformers

2 Background and Related Work

- Enable LLMs to learn from massive datasets
 - Using pre-trained model weights as a weight initialization for new tasks.

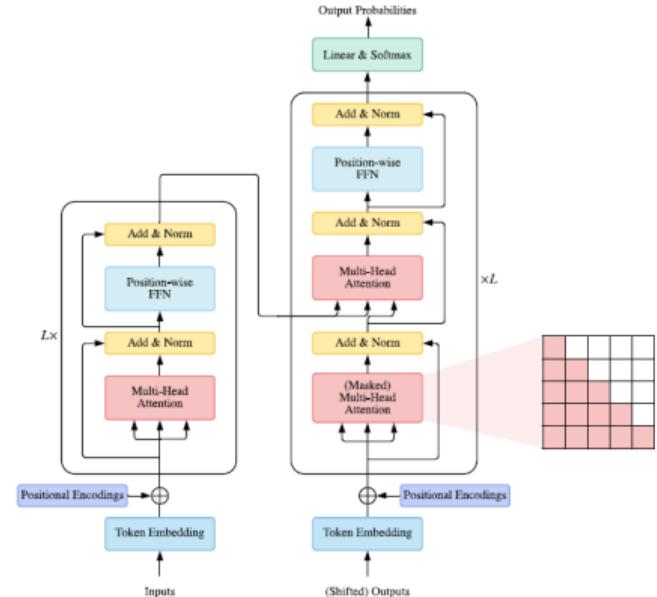




Sequence Modelling and Transformers

2 Background and Related Work

- Enable LLMs to learn from massive datasets
 - Using pre-trained model weights as a weight initialization for new tasks.
- *Evolution through Large Models (Lehman et al., 2022)*

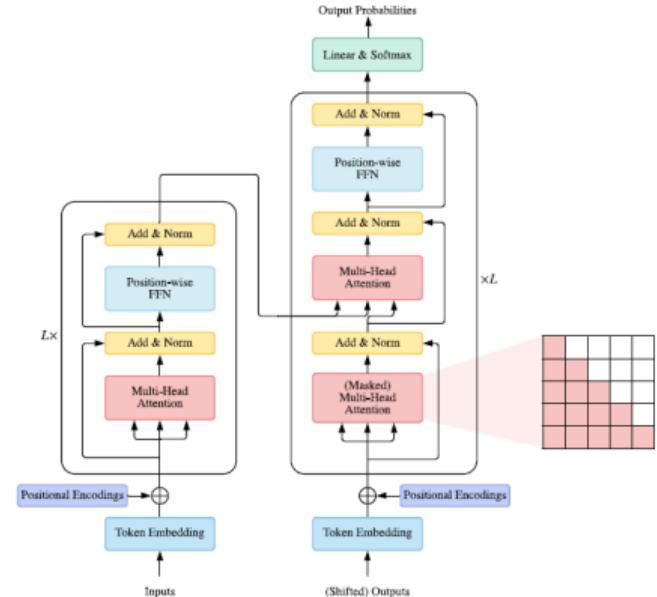




Sequence Modelling and Transformers

2 Background and Related Work

- Enable LLMs to learn from massive datasets
 - Using pre-trained model weights as a weight initialization for new tasks.
- *Evolution through Large Models (Lehman et al., 2022)*
 - LLM diff model
 - used as a "mutation operator"





Sequence Modelling and Transformers

2 Background and Related Work

- Enable LLMs to learn from massive datasets
 - Using pre-trained model weights as a weight initialization for new tasks.
- *Evolution through Large Models (Lehman et al., 2022)*
 - LLM diff model
 - used as a "mutation operator"
- Produce incredibly diverse mutations, vary increasingly over the course of the GA!

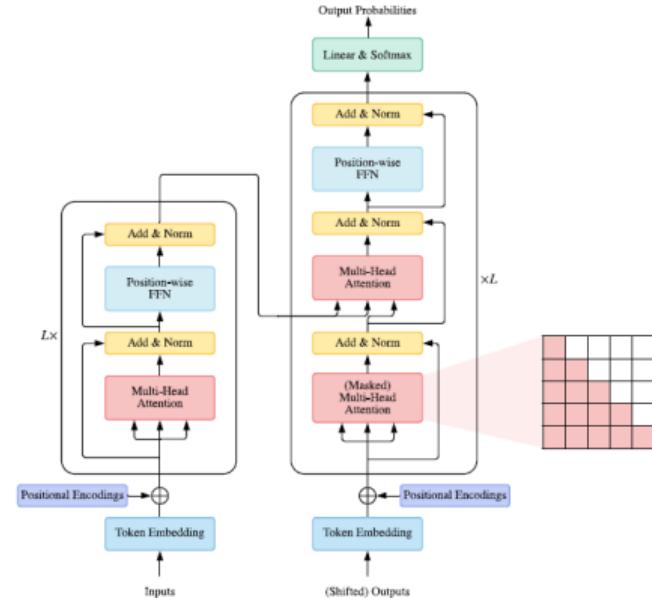




Table of Contents

3 Open-Ended Level Generation through LLMs

- ▶ Introduction
- ▶ Background and Related Work
- ▶ Open-Ended Level Generation through LLMs
- ▶ Experiments and Results
- ▶ Conclusion



Level Representation

3 Open-Ended Level Generation through LLMs

- Use the levels provided in the Video Game Level Corpus ([Summerville et al., 2016](#))



Level Representation

3 Open-Ended Level Generation through LLMs

- Use the levels provided in the Video Game Level Corpus ([Summerville et al., 2016](#))
- Utilize a set of path-annotated levels, taken from Super Mario Bros. and Super Mario Bros.: The Lost Levels (in total 37 levels)



Level Representation

3 Open-Ended Level Generation through LLMs

- Use the levels provided in the Video Game Level Corpus ([Summerville et al., 2016](#))
- Utilize a set of path-annotated levels, taken from Super Mario Bros. and Super Mario Bros.: The Lost Levels (in total 37 levels)
 - Stitch selected levels together, make one giant level!



Level Representation

3 Open-Ended Level Generation through LLMs

- Use the levels provided in the Video Game Level Corpus ([Summerville et al., 2016](#))
- Utilize a set of path-annotated levels, taken from Super Mario Bros. and Super Mario Bros.: The Lost Levels (in total 37 levels)
 - Stitch selected levels together, make one giant level!
- Tokenize string representation into discrete values using Byte Pair Encoding.



Level Representation

3 Open-Ended Level Generation through LLMs

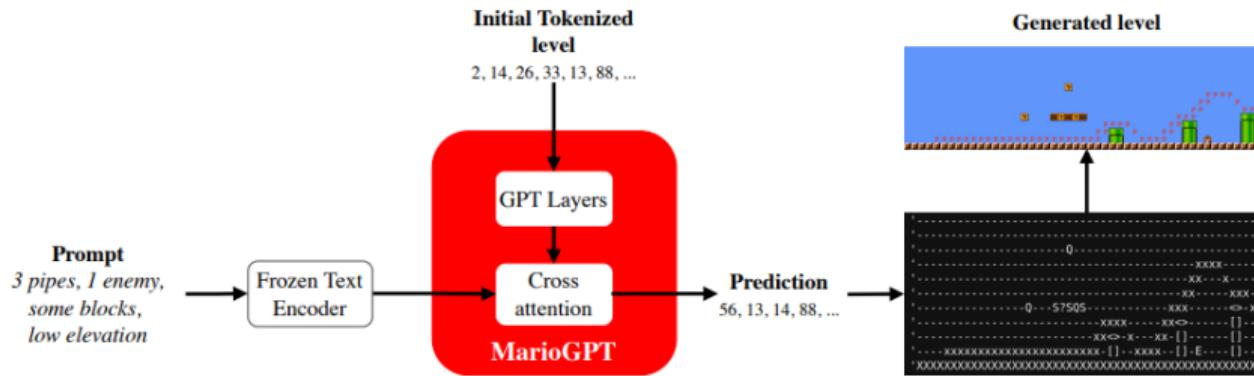
Tile Type	Symbol	Visualization
Empty	-	
Unbreakable	X	
Breakable	S	
Question Block	? / Q	
Coin	o	
Enemy	E	
Left pipe top	<	
Right pipe top	>	
Left pipe lower	[
Right pipe lower]	
Cannon Top	B	
Cannon Body	b	
Path	x	

Unique Mario tiles ([Sudhakaran et al., 2023](#))



MarioGPT Model

3 Open-Ended Level Generation through LLMs

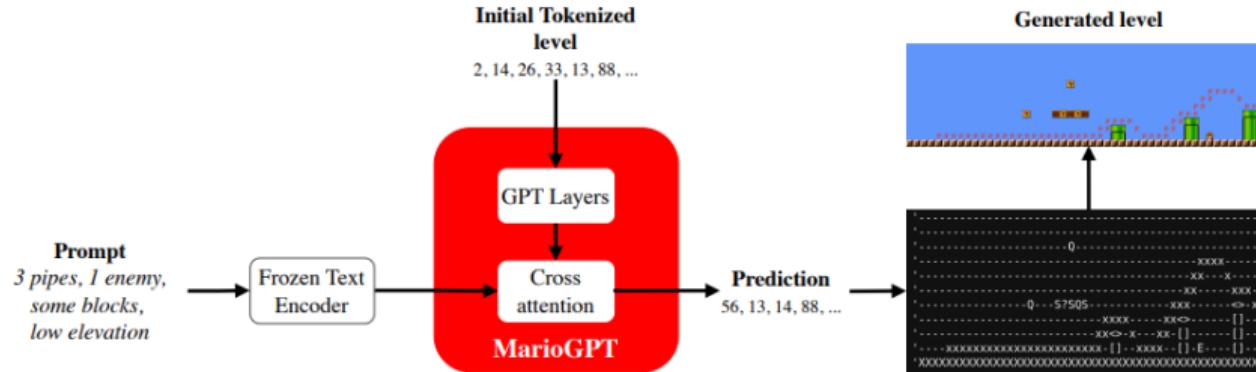


- DistilGPT2 (Li et al., 2021) is MarioGPT’s backbone.



MarioGPT Model

3 Open-Ended Level Generation through LLMs

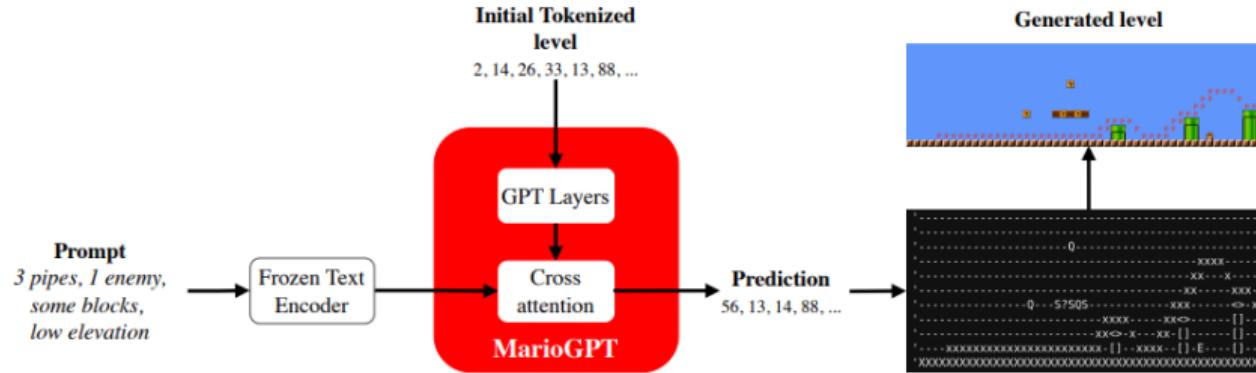


- DistilGPT2 (Li et al., 2021) is MarioGPT's backbone.
- Cross attention layer takes prompt information into account.



MarioGPT Model

3 Open-Ended Level Generation through LLMs

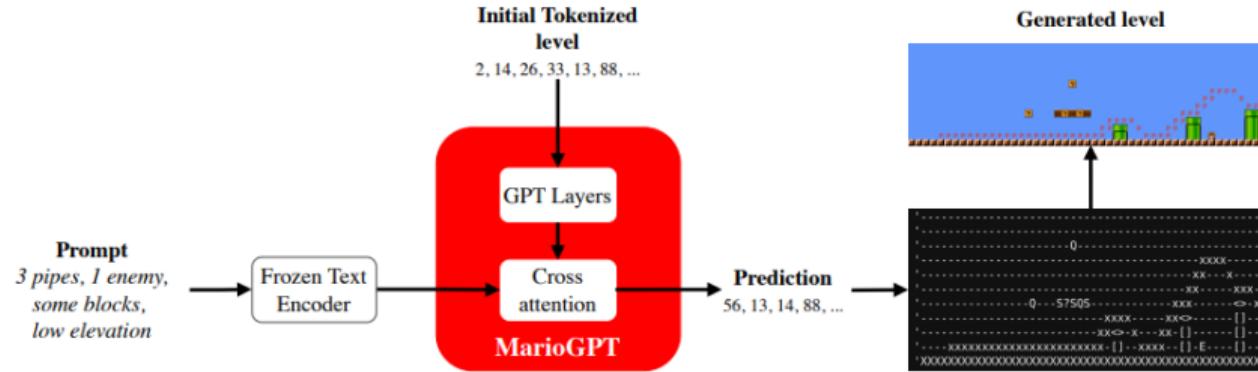


- DistilGPT2 (Li et al., 2021) is MarioGPT's backbone.
- Cross attention layer takes prompt information into account.
- 96M parameters of DistilGPT2 (86M) and Cross attention weights (10M).



MarioGPT Model

3 Open-Ended Level Generation through LLMs



- Prompts are encoded through BART ([Lewis et al., 2019](#)).
- Prompt's feature vector is combined with the actual level sequence in the cross attention layers.



MarioGPT Model

3 Open-Ended Level Generation through LLMs

- Prompts:
 - $\{ \text{no, little, some, many, [0-1000]} \} \text{ pipes}$
 - $\{ \text{no, little, some, many, [0-1000]} \} \text{ enemies}$
 - $\{ \text{little, some, many, [0-1000]} \} \text{ blocks}$
 - $\{ \text{low, high} \} \text{ elevation}$



MarioGPT Model

3 Open-Ended Level Generation through LLMs

- Prompts:
 - { no, little, some, many, [0-1000] } pipes
 - { no, little, some, many, [0-1000] } enemies
 - { little, some, many, [0-1000] } blocks
 - { low, high } elevation
- Example:
 - "no pipes, many enemies, low elevation"
 - "many pipes, many enemies, many blocks"



MarioGPT Model

3 Open-Ended Level Generation through LLMs

tile	no	little	some	many
pipes	0	1	2	5
enemies	0	1	3	7
blocks	0	50	75	176

Prompt Quantiles and corresponding
counts within a 50 column window



MarioGPT Model

3 Open-Ended Level Generation through LLMs

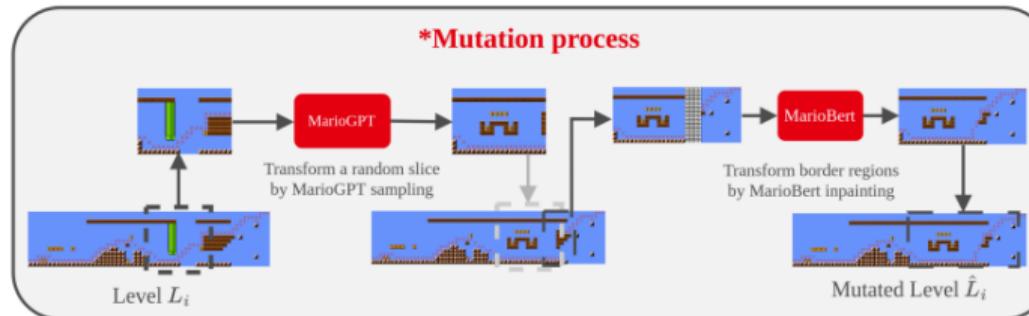
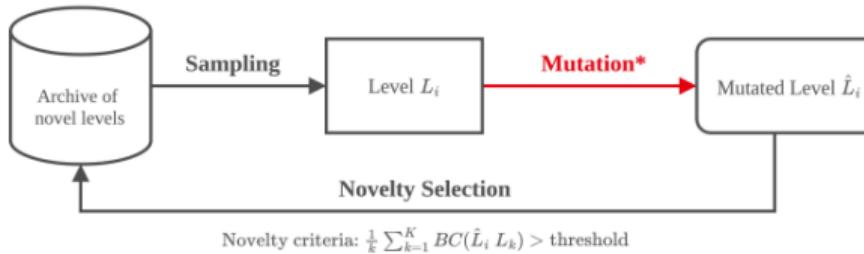




Table of Contents

4 Experiments and Results

- ▶ Introduction
- ▶ Background and Related Work
- ▶ Open-Ended Level Generation through LLMs
- ▶ Experiments and Results
- ▶ Conclusion



Table of Contents

5 Conclusion

- ▶ Introduction
- ▶ Background and Related Work
- ▶ Open-Ended Level Generation through LLMs
- ▶ Experiments and Results
- ▶ Conclusion



MarioGPT: Open-Ended Text2Level Generation through Large Language Models

Thank you for listening!

Any questions?